

Adaptive Filter Reordering for JIT-Compiled Queries in PostgreSQL

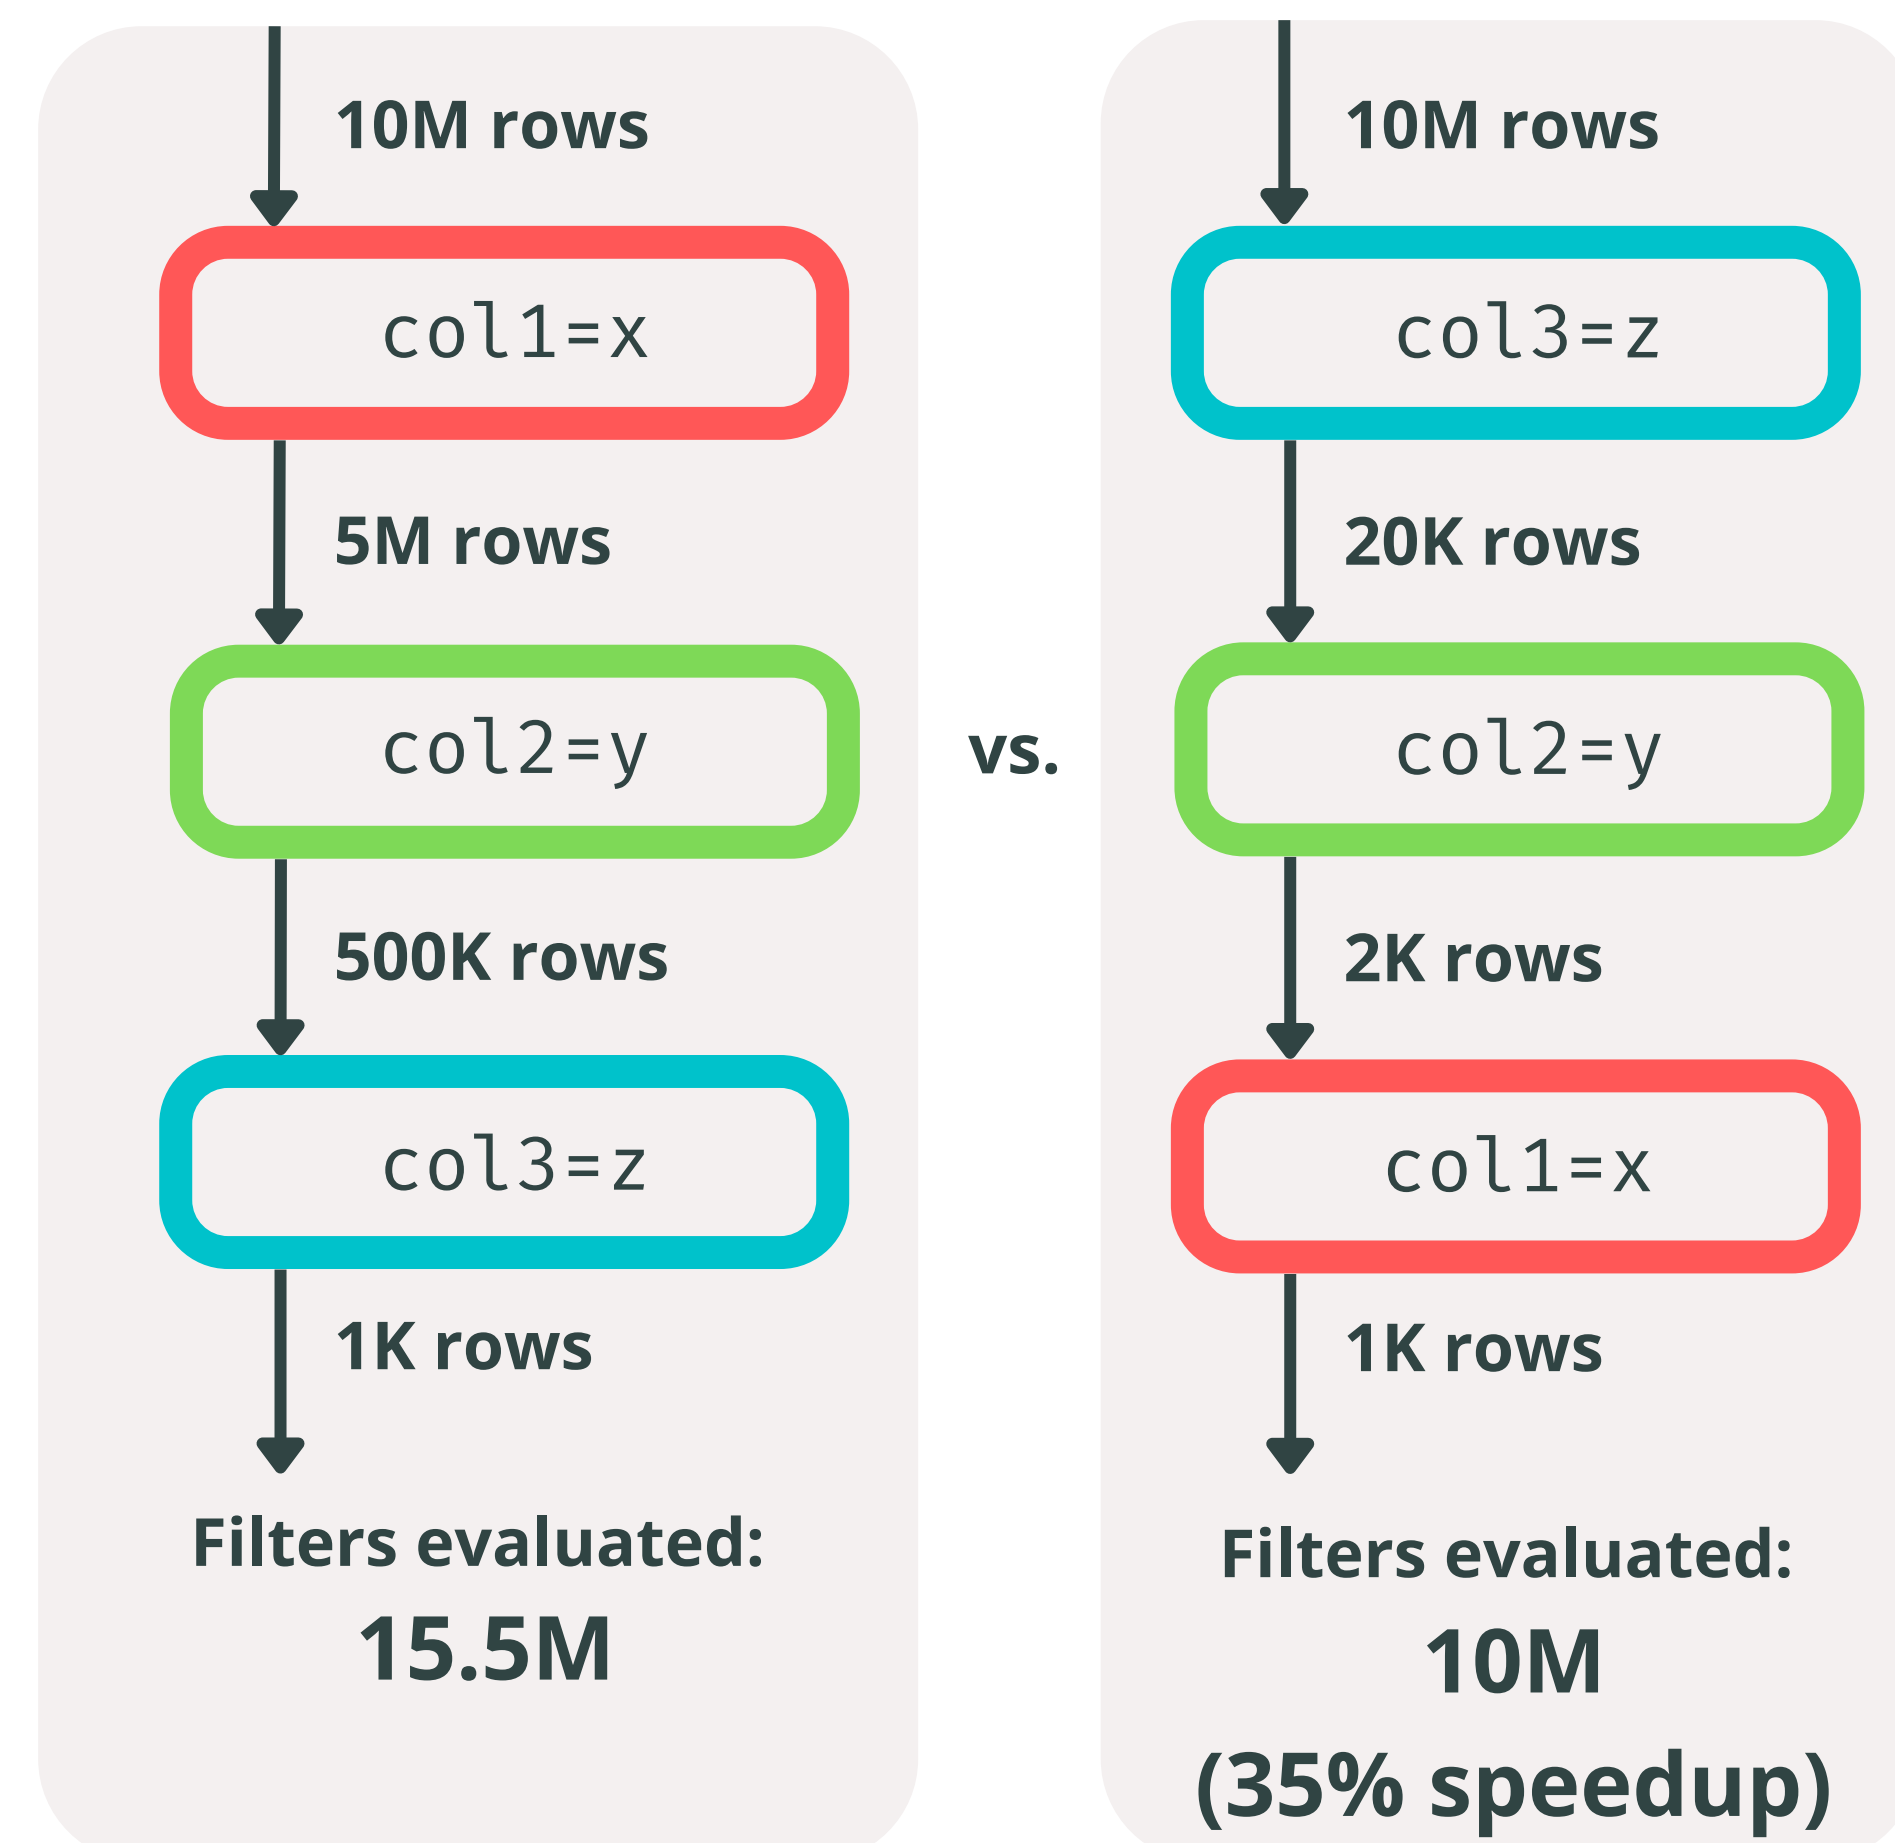
Kai Franz (kfranz@andrew.cmu.edu)

Advised by Todd C. Mowry and Andrew Pavlo

Motivation

- Filtered SELECT database queries
 - Conjunction of several filter clauses
 - e.g. **WHERE col1=x**
 - Evaluation order affects query latency
- Optimal ordering is data-dependent
 - Often unknown until runtime

```
SELECT * FROM foo
WHERE col1=x
AND col2=y
AND col3=z
```

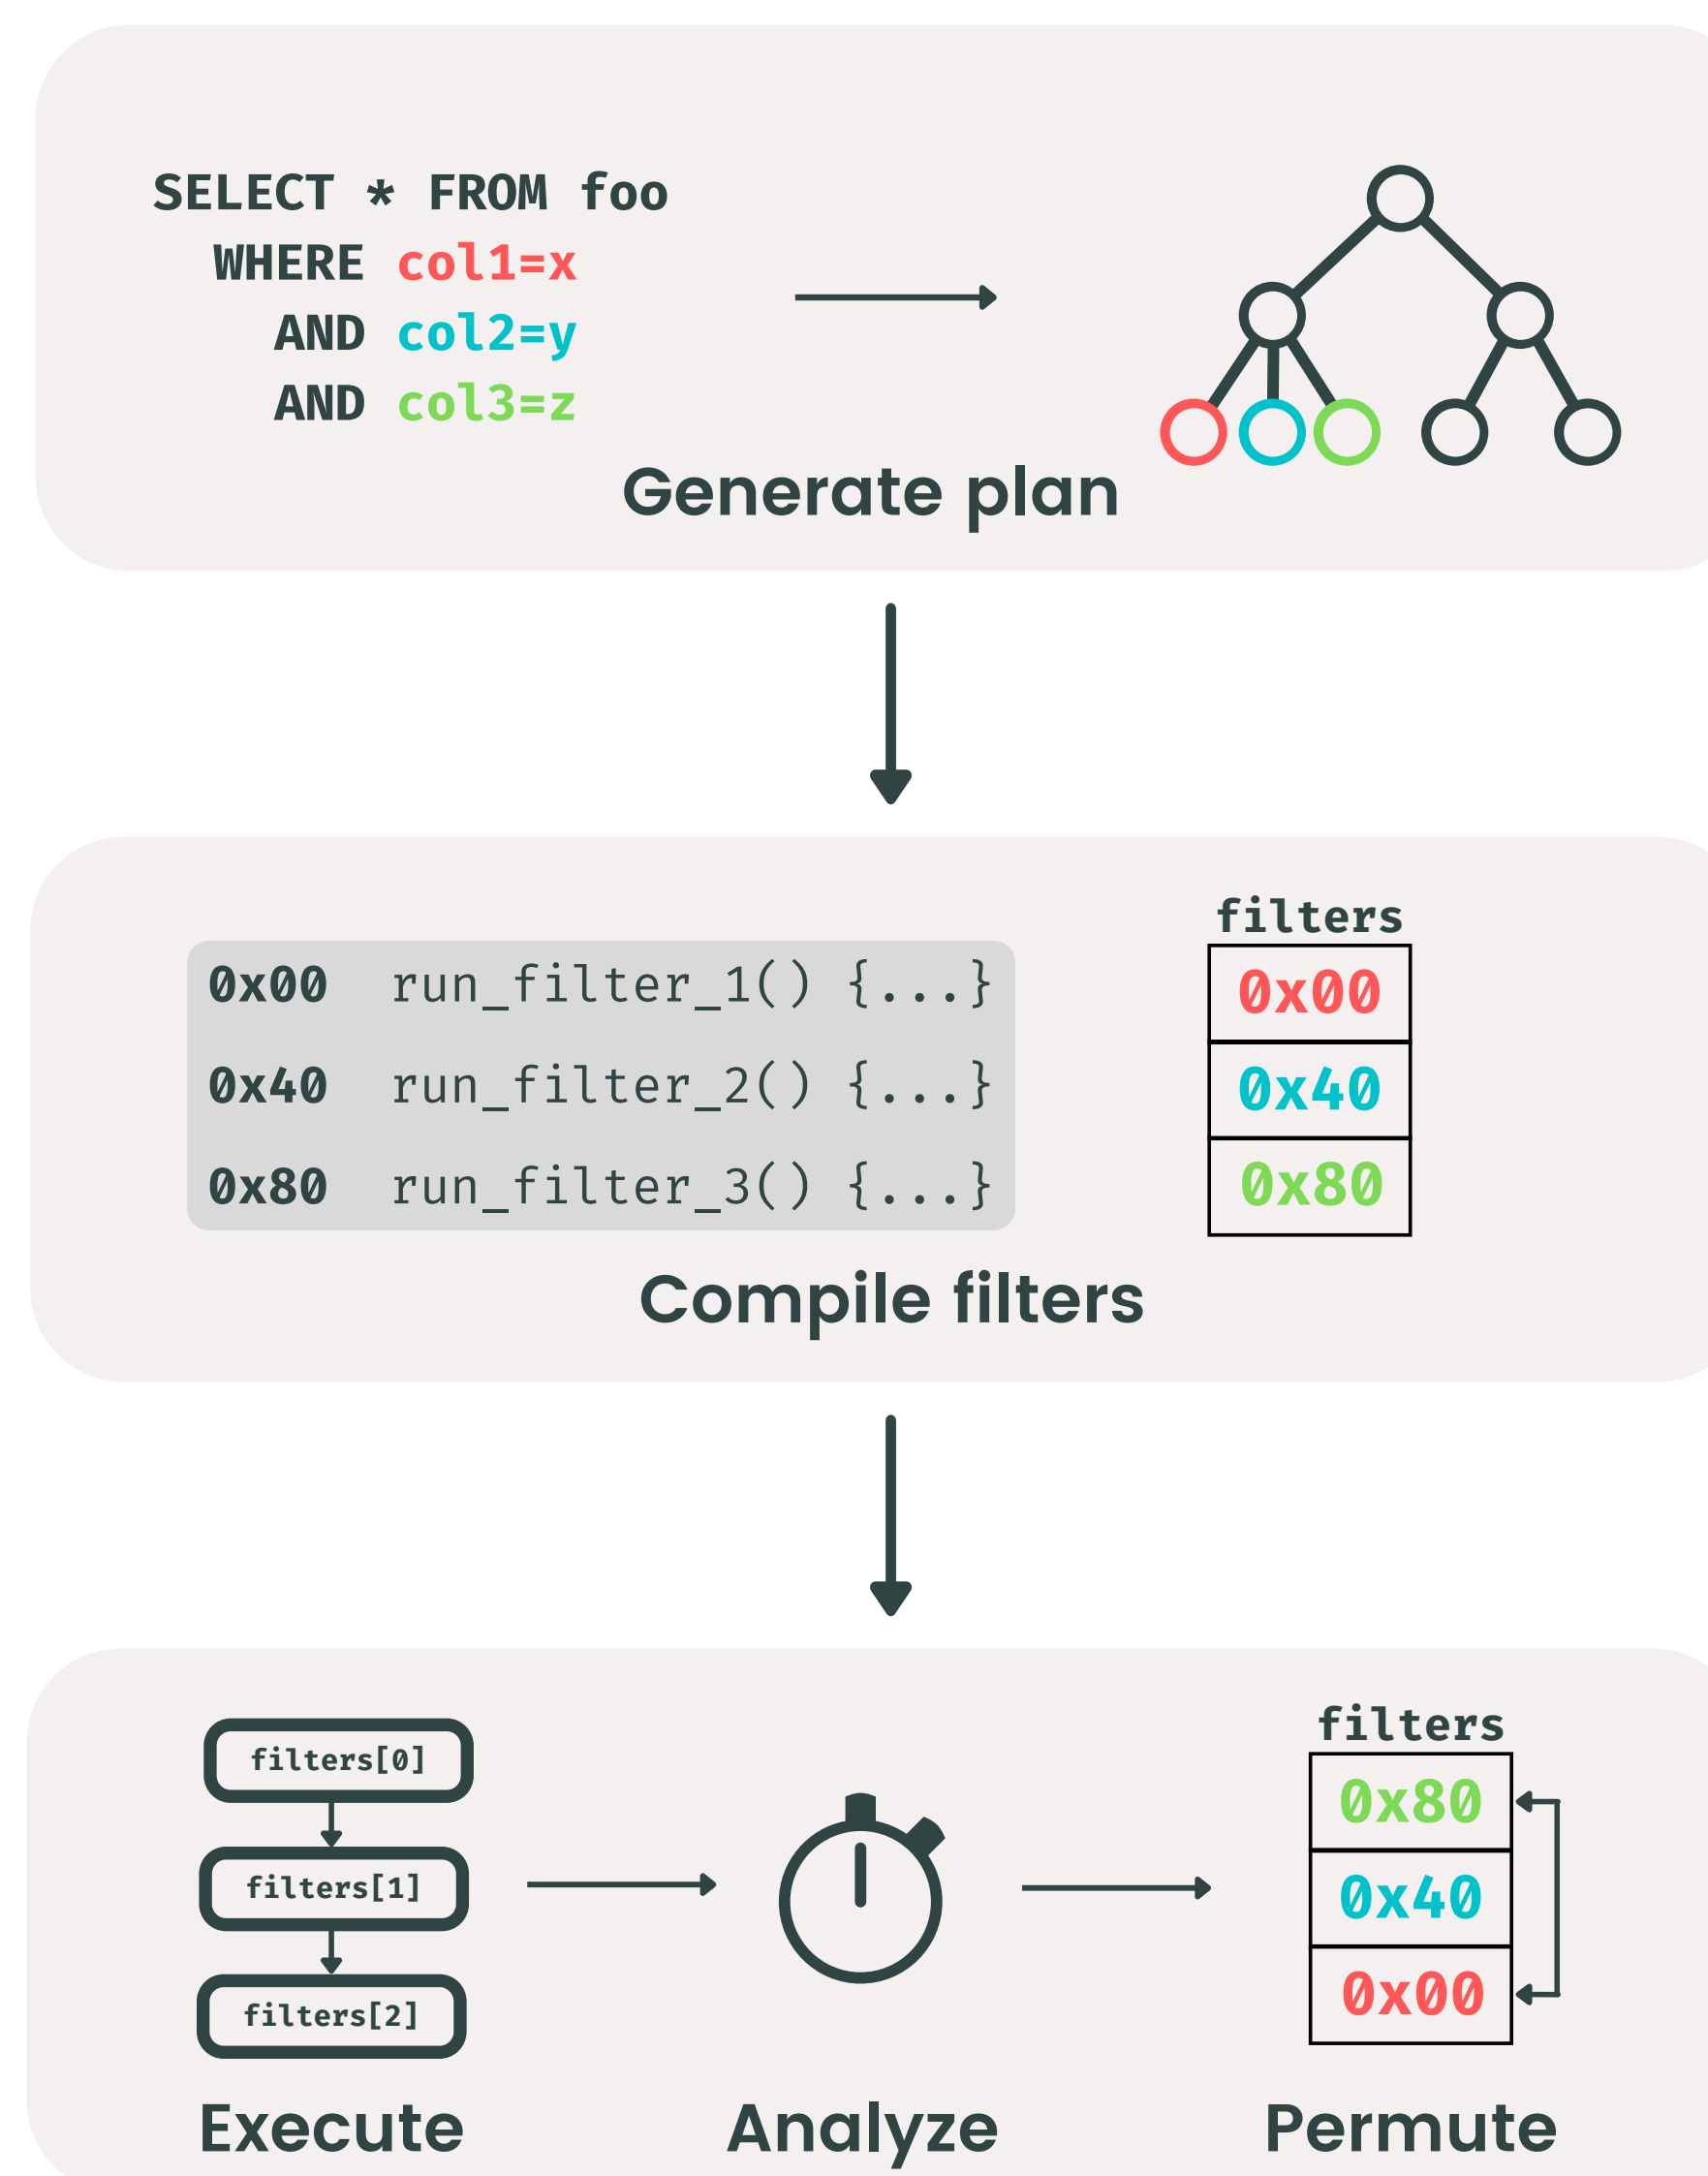


References

1. Menon, Prashanth, et al. "Permutable Compiled Queries." Proceedings of the VLDB Endowment, vol. 14, no. 2, 2020, pp. 101–113., <https://doi.org/10.14778/3425879.3425882>.

Permutable Compiled Queries¹

- Compile each filter as a separate function
- Keep table of pointers to each filter
- Periodically gather statistics for each filter
 - Execution time
 - Selectivity
- Reorder table entries based on statistics
- Allows for runtime permutation
- Avoids expensive recompilation




Implementation

- Modify PostgreSQL's JIT compiler
 - Subexpressions of qual expressions
 - Build separate functions in LLVM
- First time expression is evaluated
 - Compile LLVM module
 - Initialize filter table
- Modify PostgreSQL executor
 - Add filter manager that executes filters as usual
 - Collects runtime statistics
 - Permutes filter table

Evaluation

- Microbenchmark: in-memory table with 3 64-bit integer columns, 10M rows
- Query: 3 filters with identical execution time
 - Varying selectivities
 - 88%, 92%, 0.5%
- Adaptive reordering finds optimal filter ordering that static optimizer misses



Future Work

- More runtime optimizations
 - Adaptive Aggregations
 - Adaptive Joins
- Reuse of compiled query plans